Consider the following training data, based on which we train a support vector machine (SVM).

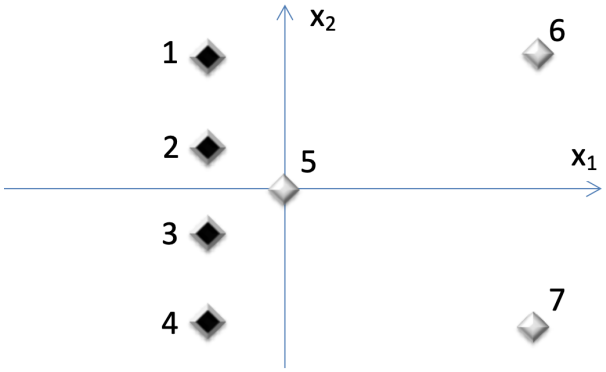

(a) In the figure above, draw the decision boundary obtained by the SVM, as well as the lines defining the margin of separation.

(b) What is the classification accuracy of the training samples?

(c) The removal of which sample would move the decision boundary?

(d) If we use a leave-one-sample-out (LOSO) cross-validation, what would be the average classification accuracy based on all folds?
Note: During LOSO we have as many folds as the training data. In each fold, we assume one sample in the test and the remaining in the train.